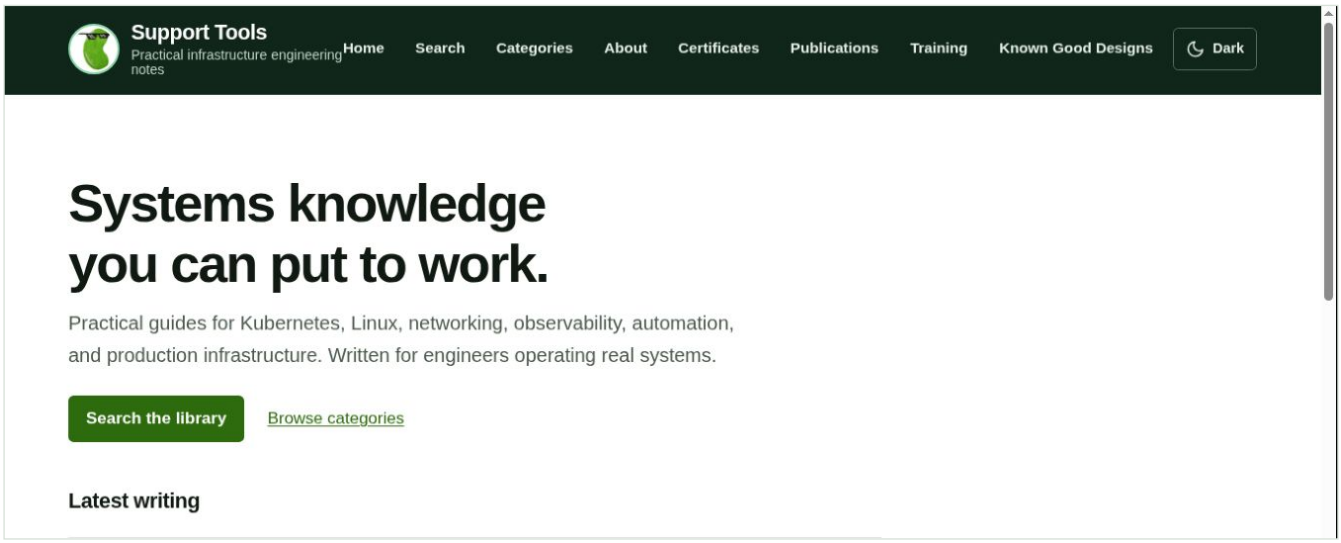
01 / ASSESSMENT

A strong reading foundation. A harder path to the right answer.

Review of https://support.tools/ for Matthew Mattox. Findings reflect the public site observed on 20 October 2026.

The site has a coherent green visual identity, clear technical positioning, and a comfortable reading layout. The most valuable next step is to improve how engineers find, assess, and use the right guide. Four issues deserve the first round of work: low-contrast search matches, fragmented categories, empty training paths, and incomplete navigation within design guides. [S1-S5]



Evidence 1. The homepage communicates the subject and audience clearly. Search and categories are already prominent; the crowded header is a separate issue.

Preserve what already works

Keep the text-first layout, generous whitespace, restrained color palette, and constrained article column. The live article sample has a working sticky table of contents with active-section feedback. Code blocks provide a Copy button and visible Copied feedback. Both light and dark themes are cohesive. Search returns useful matches and can link directly to article sections. [S1, S2, S6]

Review basis and limits

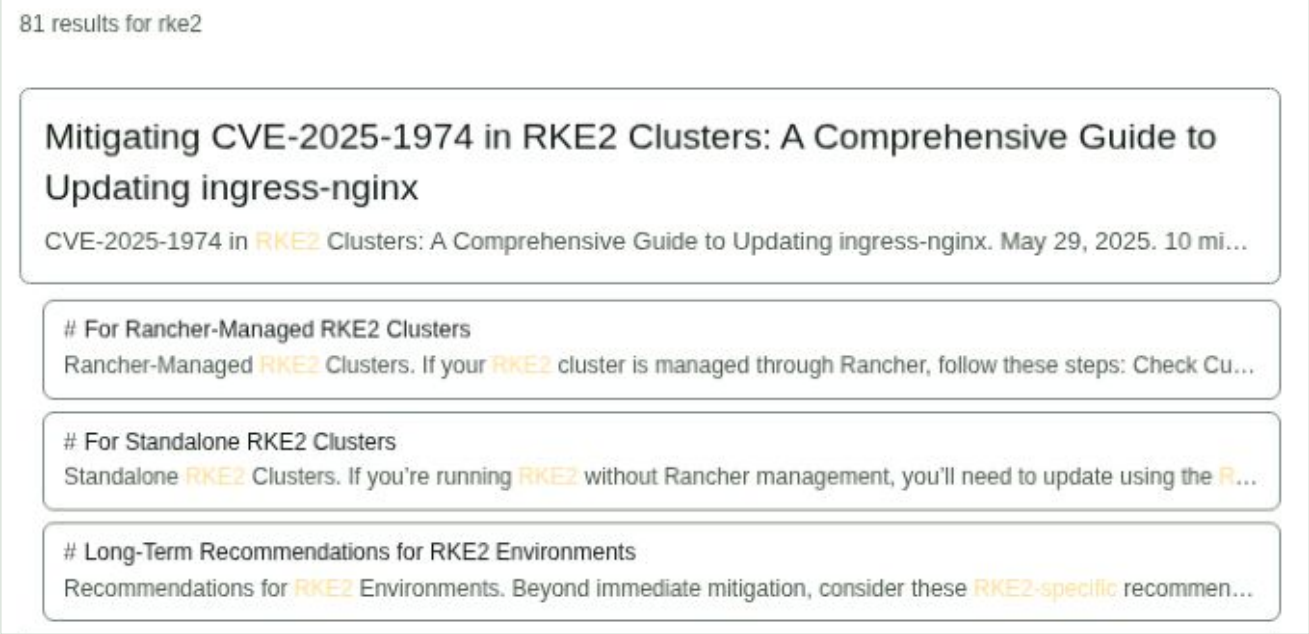
The assessment combines live desktop interaction at 1363 x 936, public-page inspection, and targeted HTML/CSS checks. It covers the homepage, search, categories, articles, training, design guides, and author pages. Mobile findings are based on responsive CSS only; no mobile viewport or physical-device task test was performed. No analytics, field performance data, Lighthouse run, full accessibility audit, or technical validation of published commands was included. Recommendations and expected benefits are judgments, not measured conversion improvements.

02 / P1 - FIX FIRST

Make search matches readable

A localized contrast defect affects the words users are actively looking for. [S2, S10]

In the live light-theme search, matched terms use pale yellow text (#ffe29a) on a transparent background over white. The measured contrast is **1.27:1**. In the dark-theme CSS, matches use #745a20 against #131c16, giving a calculated **2.68:1**; that dark search state was not measured live. Both are below the 4.5:1 WCAG minimum for normal text. This finding concerns search matches; sampled base body and muted text had strong contrast.



Evidence 2. Pale matched words become difficult to distinguish in compact result snippets. The RKE2 query returned 81 results in the live session.

Small change, immediate benefit

The search wrapper sets `--pf-mark` to `var(--color-mark)`, but Pagefind uses the variable as a foreground color. Pointing it to the existing accent color is a minimal candidate fix:

Proposed CSS change

```
.search-interface {
  --pf-mark: var(--color-accent);
}
```

With the sampled theme tokens, that candidate produces approximately **6.51:1** in light mode and **9.46:1** in dark mode. Verify the rendered result in both themes after implementation. Alternatively, use a contrasting highlight background and a separately defined dark foreground. Do not rely on a background-highlight color as the text color.

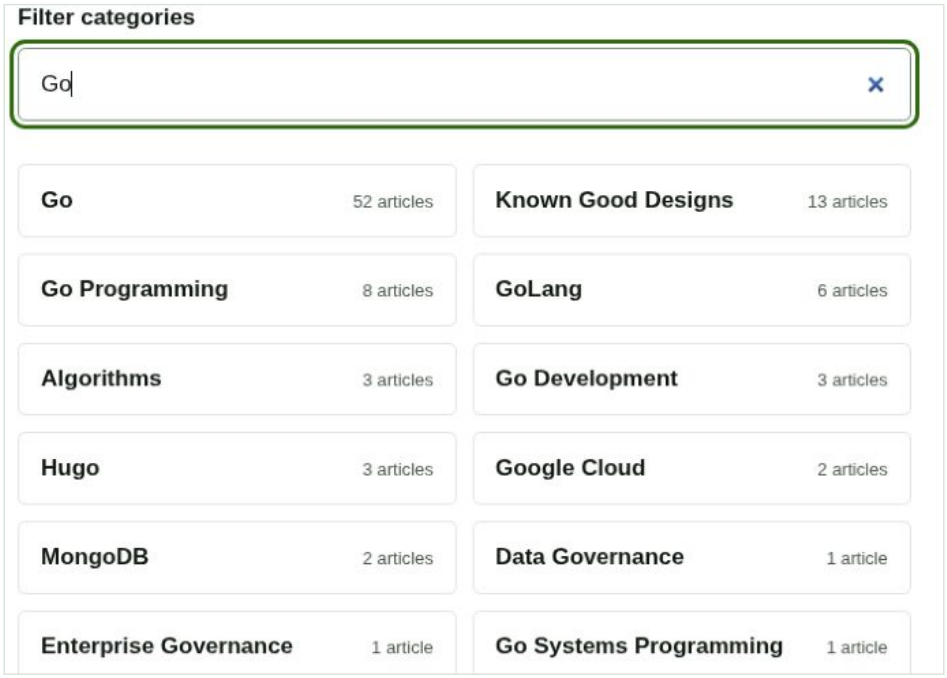
Acceptance: every matched term remains at least 4.5:1 against its actual background in light and dark themes, including focused or hovered result states. Preserve normal result text and section links.

03 / P1 - FIX FIRST

Replace category sprawl with clear topic paths

The current directory contains 445 category entries. [S3]

The directory exposes authoring vocabulary directly to readers. Go has 52 articles, while Go Programming has 8, GoLang 6, Go Development 3, and Go Systems Programming 1. These counts may overlap and should not be added. AI/ML appears twice with different counts but the same destination; SSL/TLS has the same duplicate-label pattern. Similar fragmentation affects backups, databases, and site reliability.



Evidence 3. Filtering for Go reveals several competing Go collections. Other matches, such as Hugo or Known Good Designs, follow the existing substring rule; the filter itself is not broken.

A manageable topic model

Create 6-10 canonical hubs around how readers look for help: Kubernetes and Rancher; Linux; Networking; Security; Observability; Automation and Go; Infrastructure; and Reference Designs are sensible starting candidates. Validate names against the actual library. Each hub should offer a few curated starting points, common troubleshooting tasks, and then the complete article list.

Keep granular tags as secondary facets. Normalize aliases such as GoLang to Go, merge duplicate labels deliberately, and preserve old URLs through redirects or aliases. Recalculate counts from distinct articles. Improve filtering with exact and word-prefix matches first, while retaining broader matches as secondary results.

Acceptance: no duplicate visible topic labels point to the same destination; one canonical collection exists for each agreed subject; old links still resolve; readers can reach a relevant guide without choosing between synonymous categories.

04 / P1 + P2

Finish the discovery journey

Search works, but result context, persistent queries, and training destinations need attention. [S2, S4]

P1 - Remove empty training destinations

Training links to Fleet and RKE, yet both pages contain only short introductory copy and no lessons or resource links. The Fleet page was also confirmed in the live browser. Selecting a subject ends the learning journey instead of starting it. Publish a useful lesson list, connect existing relevant articles, or label the collection as forthcoming before the reader enters it.

A useful training card should communicate subject, scope, prerequisites, approximate commitment, and the next action. A small real collection is preferable to an unfulfilled promise. Acceptance is straightforward: every advertised training path offers a lesson, resource, or explicitly stated availability status.

P2 - Give results enough context to choose

Results provide useful article-section links, but snippets are only one line, use 12-13 px text, and often spend their space repeating title, date, or tags before useful context appears. Allow two or three readable lines, emphasize the matching passage, and distinguish a main article result from a section result. Keep deep links: they are especially useful during troubleshooting.

Review Pagefind indexing boundaries so result excerpts prioritize article content. Apply body indexing consistently across templates; selectively ignore repeated navigation and metadata instead of accidentally excluding useful pages. Follow Pagefind's indexing guidance and verify representative queries before and after the change. [S11]

P2 - Persist the query in the URL

The tested query RKE2 returned 81 results; CrashLoopBackOff returned 8 and included a useful exact result. Returning from an article with browser Back restored search in the observed session. However, the URL stayed at /search/ and reloading cleared the query. Store the query in a parameter such as `?q=rke2` and restore it when the page loads.

> **Acceptance for search**
> A bookmarked or shared query reproduces the search after reload. Main result and section links remain usable. Snippets provide enough context to distinguish similar guides. Add a deliberate test for true zero results and recovery suggestions; a genuine zero-result state was not verified in this review.

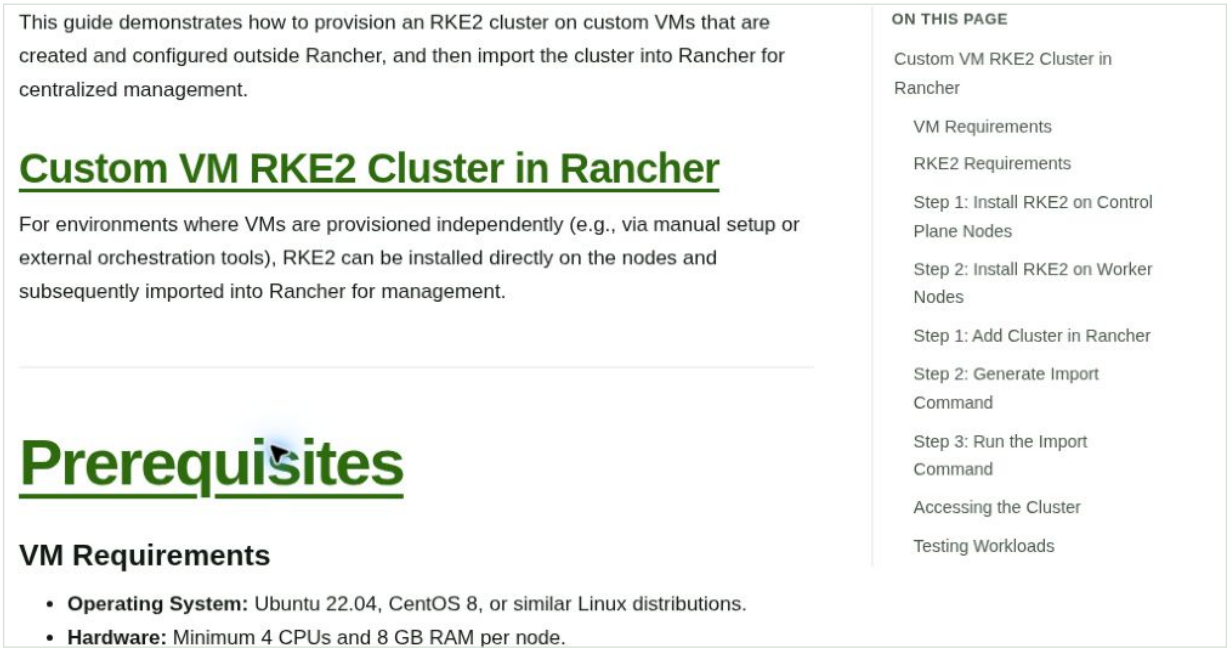
Measure relevant clicks and query recovery rather than result count alone. More results do not necessarily make an operational answer easier to find.

05 / P1 - FIX FIRST

Repair the design-guide outline

Heading structure affects visual hierarchy, in-page navigation, and assistive navigation at the same time. [S5]

The custom-VM RKE2 guide contains **seven H1 elements**, including the page title. Its major sections - Prerequisites, Setting Up RKE2, Importing the Cluster, Testing and Validation, Considerations, and References - do not appear in the generated table of contents. The visible Prerequisites heading competes with the page title. Readers instead see lower-level requirements and repeated step labels without their full context.



Evidence 4. Prerequisites is visually oversized and absent from the adjacent table of contents. The outline jumps directly to lower-level requirements and steps.

One structural fix across the template

Use H1 for the page title, H2 for major sections, and H3 for steps. Ensure the table of contents includes every major section in the same order as the article. Give repeated step titles enough context to stand alone, particularly on the small-RKE2 guide where Create the Namespace and Helm Chart Deployment appear under different deployments.

The custom-VM breadcrumb labeled Known Good Designs currently goes to the category archive instead of the curated design collection. Point it to `/known-good-designs/` so moving up restores the collection context. Add lightweight source checks for extra H1s, skipped heading levels, and expected collection parents.

Acceptance: one H1 per article; a complete and correctly nested outline; every TOC link lands on its named section; the collection breadcrumb returns to the curated hub. Retain the working sticky TOC and active-section feedback on article templates.

06 / P2 - NEXT PASS

Show when a guide applies

For an engineering reference, confidence depends on applicability as well as appearance. [S5-S9]

Make the Known Good promise assessable

The design hub describes validated configurations. Two sampled RKE2 guides show a publication date of December 17, 2024 but no visible last-validation date or tested-version matrix. That does not prove the designs were untested. It does mean the reader cannot determine what the validation covers or whether the guide matches the environment they operate today.

Applicability panel	What the reader should learn
Tested versions	Rancher, Kubernetes/RKE2, operating system, and relevant chart versions.
Last validated	Date of the most recent documented verification, separate from publication date.
Environment and scope	Lab or production scope, topology, prerequisites, and known limitations.
Ownership and feedback	Named author/reviewer and a clear way to report a broken or outdated step.

Populate those fields from real evidence. Where validation is historical or unknown, state that plainly. A technical review of the underlying instructions is separate work; this review does not establish that commands are correct or current.

Reduce metadata that repeats itself

Articles show category and tag rows that repeat topics. Recent titles often combine product names with broad enterprise or production framing. Similar NGINX tuning titles make the distinction between a quick reference and a deep guide difficult to see. Present one primary topic and a useful content type; place secondary tags farther down or behind disclosure.

Use titles that distinguish the actual task or outcome. For example, a troubleshooting title should name the symptom and platform; a reference design should name its topology. Mark quick reference, tutorial, and architecture guide explicitly. Avoid adding production or validated labels unless the supporting content establishes them.

Make the next step relevant

Prefer related technical guides or the next lesson over chronological neighbors. The sampled Pixie article links onward to unrelated subjects by date. Static author pages inherit article pagination: About ends with Older article: Certificates. Use separate templates for static pages and learning series. The About page also contains a CV sentence with no actual link; connect it to a real destination or remove the promise.

07 / P2 - NEXT PASS

Give navigation room and readers a starting point

Refine the existing identity and layout; the core reading experience deserves to stay. [S1, S3, S6]

Resolve the desktop header collision

At the observed desktop width, the brand tagline touches Home. The header switches to a single row at 64 rem while navigation is nonwrapping and allows visible overflow. The available width does not reliably accommodate the brand, eight links, and theme control. Define an explicit gap, protect the brand from shrinking into the navigation, and keep a two-row layout or menu until the full set genuinely fits.

Group reader tasks around Search, Topics, Training, and Reference Designs. Move author credentials and publications into an About grouping where appropriate. Rename Certificates to Certifications so it is clearly about professional credentials rather than TLS material. Keep the theme switch easy to find and preserve keyboard access when navigation collapses.

Make the homepage a useful library entrance

The current homepage already explains the audience and offers search/category actions. Its primary content then becomes a latest-post feed across 115 pages. Add a working search input and a few curated operational paths before Latest Writing. This reduces the effort for someone arriving with a fault, deployment task, or specific technology in mind.

> **Suggested homepage order**
> Clear purpose and search field; a compact set of common tasks or topic hubs; featured reference designs or starting guides; latest writing; concise author credibility and contact. These are recommendations, not a requirement to add every section.

Use existing content to populate those paths. Keep the comfortable whitespace, article line length, plain language, and restrained palette. Screenshots or diagrams should explain technical material when useful; decorative stock imagery would consume space without helping readers solve an operational problem.

Mobile and accessibility follow-up

The keyboard-tested skip link becomes visible with a focus outline and Enter moves to main content. CSS and markup also include 44 px navigation targets, reduced-motion handling, labeled search, and status regions. Those checks do not establish complete accessibility. Test the actual tasks at 320, 390, and 768 px, then with keyboard navigation and zoom: open the menu, search, choose a section, read or scroll code, and return to results.

Check for page-wide horizontal scrolling, clipped focus, crowded controls, and obscured anchor targets. Verify the mobile TOC and theme switch in context. Mobile completion and field performance remain unmeasured; this report makes no numerical score or pass claim for them.

08 / DELIVERY PLAN

Fix failure points first

P1 affects successful use of core content. P2 improves discovery, confidence, and polish. S/M are relative implementation sizes, not time estimates.

Work item	Priority / effort	Acceptance condition
Search match contrast	P1 / S	All matched terms meet 4.5:1 in both themes using actual rendered backgrounds.
Canonical topic model	P1 / M	Duplicate labels and synonyms are normalized; distinct counts and legacy URLs remain correct.
Fleet and RKE training paths	P1 / S-M	Each advertised path contains useful resources or clearly states that it is forthcoming.
Design headings and breadcrumb	P1 / S-M	One H1; complete nested TOC; collection breadcrumb restores the design hub.
Search snippets and query URL	P2 / M	Readable contextual excerpts; query survives share, bookmark, and reload; section links remain useful.
Header and navigation grouping	P2 / S-M	Brand and navigation fit at transition widths; menu remains keyboard accessible.
Homepage task paths	P2 / M	Readers can search immediately and enter curated paths built from existing useful material.
Guide applicability metadata	P2 / M	Sampled design guides show documented versions, scope, validation date, and a feedback route.
Related content and static pages	P2 / S-M	Next steps match the subject; About has a working CV link and no article chronology.

Practical order of work

Ship the contrast fix, query persistence, and header spacing as early quick fixes; resolve empty training destinations alongside them. Repair guide heading structure and collection context next. Normalize taxonomy before expanding homepage topic navigation, so new routes use stable vocabulary. Then refine search excerpts and introduce the applicability panel across maintained guides.

The taxonomy and applicability work require editorial decisions as well as template changes. Agree on canonical terms and evidence requirements before bulk updates. A small reviewed set of high-use guides is a sensible starting point for validation metadata.

Release checks: repeat RKE2 and CrashLoopBackOff searches; exercise share/reload and section links; inspect headings and breadcrumb destinations; review both themes; and perform the proposed mobile tasks. These checks address specific remaining risks without treating the review as a full audit of every article.

09 / VERIFICATION & SOURCES

Measure whether answers are easier to reach

Establish a baseline before claiming improvement. The measures below are proposed, not observed results.

Measure	What to look for
Search to relevant click	Whether readers choose an appropriate result or section for their query.
Zero or no-use search outcomes	Queries with no results, no useful click, or repeated reformulation.
Guide usefulness feedback	Whether readers completed the task and which steps were unclear or outdated.
Dead-end destination count	Advertised collections with no useful lesson, resource, or next step.
Taxonomy duplicate count	Duplicate labels and synonymous collections remaining after normalization.
Mobile task completion	Completion and friction for menu, search, section navigation, code, and return paths.

Use a combination of lightweight analytics and direct task observation. An engineer finding a copied command quickly can be successful even with a short visit. A longer session may indicate useful reading or difficulty; interpret behavior in relation to the task.

Sources and evidence

Site observations and screenshots were captured or refreshed on 20 October 2026. Public URLs below identify the reviewed surfaces. Contrast values come from live computed styles or exported CSS; interactions come from live desktop checks. No third-party traffic or performance estimates were used.

[S1] Home - https://support.tools/

[S2] Search - https://support.tools/search/

[S3] Categories - https://support.tools/categories/

[S4] Training and empty paths - https://support.tools/training/ ; https://support.tools/training/fleet/ ; https://support.tools/training/rke/

[S5] Design hub and sampled guides - https://support.tools/known-good-designs/ ; https://support.tools/known-good-designs/downstream-rke2-cluster/rke2-custom-vms/ ; https://support.tools/known-good-designs/upstream-rancher-cluster/rancher-rke2-small/

[S6] Sampled articles - https://support.tools/pixie-ebpf-observability-kubernetes-production/ ; https://support.tools/kubernetes-crashloopbackoff-example/

[S7] Kubernetes archive - https://support.tools/categories/kubernetes/

[S8] About and credentials - https://support.tools/about/ ; https://support.tools/certificates/

[S9] Publications - https://support.tools/publications/

[S10] W3C, Understanding Contrast (Minimum) - https://www.w3.org/WAI/WCAG22/Understanding/contrast-minimum.html

[S11] Pagefind indexing guidance - https://pagefind.app/docs/indexing/